

COMP421 Crib Sheet Francis Piché

Transactions

- A sequence of reads **r(x)** and writes **w(x)**
- Atomic** (all or nothing)
 - Keep *backup* of state before transaction
 - Restore to this point in case of failure
- Consistency** (preserve consistency)
- Isolation** must have serial equivalent
- Durability** must be permanent/fault tolerant

Transactions can be **aborted**

-*Global recovery*:

- Transactions committed before crash are in effect.
- Transactions aborted before crash are reversed
- Transactions active at time of crash are reversed
- Assume disk doesn't crash

Logs: are kept because holding back writes is insufficient.

Limited number of buffer frames means transactions cannot all be atomic.


- Log writing is fast since logs are *append-only*. Save on seek time.

Concurrency:


- Transactions need to run in isolation
- Must have concurrency control protocol to enforce this
- Ensure net effect of concurrent transactions is equivalent to some serial order

-**Schedules**: sequence of actions (reads/writes)

Serial Schedule




Non serial Schedule



- Unrepeatable read**: Two or more reads that give different results (another transaction changed the value in between).
- Lost Update**: A write of T1 overwritten by the write of T2.
- Dirty Read**: Read value that doesn't exist (was undone by an abort later)
- Dirty Write**: (permanent damage)
- Conflicting Operations**:
 - Same object being accessed
 - WW, WR(not RR)
- Conflict Equivalent**: Schedule is conflict equiv if:
 - Every pair of conflicting actions is ordered same way
 - Same actions of same committed transactions
- Conflict Serializable***: if:
 - Equivalent to some serial schedule with actions of schedule

Dependency Graphs:

-Not serializable if there is a cycle in the dependency graph:



-Edge formed if first operation conflicts with later one (edges are always downward in time)

-**Forming serial schedules**:

- Choose node with no incoming edges
- Put in in the schedule, delete it and all outgoing edges
- Repeat until no more nodes remaining

Locking:

-Transactions must acquire shared lock S

-Exclusive lock X before writing

-X blocks all other operations

-S blocks only writes

	S	X
S	✓	--
X	--	--

-**Phase 1**: Acquire locks when needed

-**Phase 2**: Release locks at end of transaction

-Two phase locking allows only serializable schedules

-No dirty reads/writes possible

-Transaction *cannot acquire same lock twice*


-No need to acquire S on resource if already have X for resource

-Locks are managed using **lock table**

- Entry for each resource that is locked
- Pointer to queue of locks granted
- Pointer to queue of lock requests (waiting)
- A transaction has only one lock per object
- If T has S and requests X, S is upgraded to X.
- Keep track of type of lock held
- Pointer to list of locks held by each T
- Locking/Unlocking is atomic

Submission order 2

R1(A) W2(B) W1(B) W2(A)



Deadlock:

-Deadlock is possible with two phase locking. (SEE LEFT)

Wait-For Graph:

- Like dependency graph but arrows are BACKWARDS
- Edge from Ti to Tj if Ti waits for Tj to release lock.
- Cycles mean there is deadlock
- Avoid deadlock by breaking cycles.

-Can try timeout but how long should you wait?

-Can try to request all locks at beginning of transaction (loss of concurrency)

-Optimistic concurrency control: Try transaction (no locking), if conflict, abort.

Snapshots:

- Writers make new copy
- Readers use old copy

Transactions in Java:

```
con.setAutoCommit(false);
try {
    stmt.executeUpdate("INSERT INTO Skaters " +
        "VALUES (123, 'Lilly', 18, 10)");
    stmt.executeUpdate("INSERT INTO Skaters " +
        "VALUES (345, 'Debby', 12, 10)");
    con.commit();
} catch (SQLException ex) {
    System.err.println("SQLException: " +
        ex.getMessage());
    con.rollback();
}
```

Phantoms:

- Can arise when new entries being added concurrently
- Locking can't prevent
- If inserting while doing aggregation, aggregate gets weird values

Isolation Levels:

Isolation Level/ Anomaly	Dirty Read	Unrepeatable Read	Phantom
Read Uncommitted	maybe	maybe	maybe
Read Committed	no	maybe	maybe
Repeatable Read	no	no	maybe
Serializable	no	no	no

- Read uncommitted = Read operation does not set locks, can read uncommitted writes
- Read Committed = Do not read uncommitted writes. Release read lock immediately after reading.
- Repeatable reads = standard S locking on reads
- Serializable = lock entire relation


Isolation Levels

```
TRANSACTION_READ_UNCOMMITTED
TRANSACTION_READ_COMMITTED
TRANSACTION_REPEATABLE_READ
TRANSACTION_SERIALIZABLE
```

```
con.setTransactionIsolation
(TRANSACTION_SERIALIZABLE);
```

Graph Databases: (FLEXIBLE)

- Each vertex has own properties
- Properties are K-V pair
- Can easily be extended. No pre-planning required
- Edges can have properties too (are directional)



Cypher:

TRAVERSALS:

General	DISTINCT
Math	+, -, *, /, %, ^
Comparison	=, <, <=, >, >=, IS NULL, IS NOT NULL
String comparison	STARTS WITH, ENDS WITH, CONTAINS
Boolean	AND, OR, XOR, NOT
String operators	+ (Concatenation), =~ (regex matching)

Can combine conditions by comma separating:

How to find a list of people who manages someone who mentors more than one employee ?

```
MATCH (b:Employee)-[:MANAGES]->(m:Employee)
      ,(m)-[:MENTORS]->(e1:Employee)
      ,(m)-[:MENTORS]->(e2:Employee)
WHERE e1 < e2
RETURN DISTINCT b
```

EACH EDGE IS TRAVERSED ONLY ONCE TO AVOID CYCLES

(e)-[*]->(n)	// All the way (outgoing edges)
(e)-[*..5]->(n)	// Up to a depth of 5 edges (outgoing)
(e)-[*3..]->(n)	// 3 or more edges (outgoing)
(e)-[*3..5]->(n)	// 3 to 5 edges (outgoing)
(e)-[*3..5]-(n)	// 3 to 5 edges (incoming)
(e)-[*3..5]-(n)	// 3 to 5 edges (incoming or outgoing)

SELECT * FROM Employees	MATCH(e:Employee) RETURN e;
SELECT email FROM Employees	MATCH(e:Employee) RETURN e.email;
... ORDER BY email	... RETURN e ORDER BY e.email;
... WHERE name = 'Janet'	MATCH(e:Empl {ename: 'Janet'}) RETURN e;
... WHERE deptid IS NULL	... WHERE NOT (e)-[:WORKS_IN]-() ...
	... WHERE e.job IS NULL ... (treat non-exist property as NULL)
INSERT INTO ...	CREATE (e:Empl {name: 'Jane'}-[:WORKS_IN]->(d:Depart {dname:'PR'}));
New edge b/w existing nodes:	MATCH (n1: Empl {eid: 101}), (n2: ...) CREATE (n1)-[:MANAGES]->(n2);
DELETE FROM ... (Must delete relationships)	MATCH(e: ...)-[r:WORKS_IN]->(d:Dep..) DELETE e, r, d;
Delete all edges connected to this node	DETACH DELETE e;